

Wind River® Simics® Ethernet Networking

USER'S GUIDE

4.6

<i>Revision</i>	4081
<i>Date</i>	2012-11-16

Copyright © 2010–2012 Wind River Systems, Inc.

All rights reserved. No part of this publication may be reproduced or transmitted in any form or by any means without the prior written permission of Wind River Systems, Inc.

Wind River, Simics, Tornado, and VxWorks are registered trademarks of Wind River Systems, Inc. The Wind River logo is a trademark of Wind River Systems, Inc. Any third-party trademarks referenced are the property of their respective owners. For further information regarding Wind River trademarks, please see:
www.windriver.com/company/terms/trademark.html

This product may include software licensed to Wind River by third parties. Relevant notices (if any) are provided in your product installation at the following location:
`installDir/LICENSES-THIRD-PARTY/`.

Wind River may refer to third-party documentation by listing publications or providing links to third-party Web sites for informational purposes. Wind River accepts no responsibility for the information provided in such third-party documentation.

Corporate Headquarters

Wind River
500 Wind River Way
Alameda, CA 94501-1153
U.S.A.

Toll free (U.S.A.): 800-545-WIND
Telephone: 510-748-4100
Facsimile: 510-749-2010

For additional contact information, see the Wind River Web site:
www.windriver.com

For information on how to contact Customer Support, see:
www.windriver.com/support

Contents

1	Introduction	5
2	Tutorial	6
2.1	Simple Virtual Network	6
2.2	Connect to a Real Network	7
3	Network Simulation	10
3.1	Ethernet Links	10
3.2	Ethernet VLAN Switch	11
3.3	Link Object Timing	12
3.4	IP Services	13
3.4.1	IP Based Routing	15
3.4.2	DHCP and BOOTP	16
3.4.3	DNS	16
3.4.4	TFTP	17
3.5	Observing Network Traffic	18
3.5.1	Traffic Monitoring Software	18
3.5.2	Ethernet Probe	18
3.5.3	Traffic Snooping	20
3.6	Injecting and Modifying Network Traffic	22
3.6.1	From a Network Dump	22
3.6.2	From an Ethernet Probe	22
4	Connecting to a Real Network	25
4.1	Preparing for the Examples	25
4.2	Connection Types	27
4.2.1	Port Forwarding	29
	The connect-real-network Command	30
	Example	30
	Incoming Port Forwarding	32
	Example	32
	Outgoing Port Forwarding	33
	Example	34
	NAPT	35
	Example	35

	DNS Forwarding	36
	Example	36
4.2.2	Ethernet Bridging	36
	Example	37
	Setting up TAP for Ethernet Bridging	38
4.2.3	Host Connection	40
	Example	41
	IP Forwarding	42
4.3	Accessing Host Ethernet Interfaces	42
4.3.1	TAP Access	42
	Unix	43
	Windows	43
	TAP devices with multiple users and simulations	44
4.4	Selecting Host Ethernet Interface	47
4.4.1	Unix	47
4.4.2	Windows	47
4.5	Performance	48
4.6	Troubleshooting	49
	Index	50

Chapter 1

Introduction

This guide explains how to use Simics to simulate Ethernet networks as well as how to connect the simulation to real networks. It starts with a couple of short tutorials on network simulation, then goes on describing all aspects of Ethernet networking in Simics.

Chapter 2

Tutorial

This tutorial gives a brief overview on how to simulate an Ethernet network in Simics, and how to connect a simulated network to a real Ethernet network.

2.1 Simple Virtual Network

Simics can simulate several computers in the same session, and connect them together using simulated Ethernet links.

Before reading on, launch the `firststeps-multi.simics` configuration: in the **File** menu, select **New Session from Script** and select the `mpc8641d-simple\firststeps-multi.simics` file as a start script. This configuration contains two identical MPC8641D cards, each with its own CPU. Simics will simulate both cards in the same session, keeping them synchronized.

Boot the two simulated machines by starting the simulation:

```
simics> c
```

The cards will be assigned the Internet Protocol (IP) addresses 10.10.0.40 and 10.10.0.41. At this time, trying to ping one of the machines from the other will fail: this is because there is no simulated network between the two machines yet.

The **eth-links** module provides several ready-to-use simulated network models. The simplest is the **ethernet-hub** which is a simple cloud, or hub, to which any number of ethernet devices can be connected to and talk together. Other variants include an Ethernet cable, with only two ends, and an Ethernet switch, that will act somewhat like a real Ethernet switch, sending packets to the right destination whenever it can.

```
simics> stop
simics> load-module eth-links
simics> new-ethernet-hub
Created instantiated 'ethernet_hub' component 'ethernet_hub0'
simics>
```

Next, let us connect the simulated network cards to the hub:

```

simics> connect mpc8641_simple_0.eth[0] ethernet_hub0.device0
simics> connect mpc8641_simple_1.eth[0] ethernet_hub0.device1
simics> ethernet_hub0.link.status
Status of ethernet_hub0.link [class eth-hub-link]
=====

Effective latency : 10.0 ms

Connected devices : ('mpc8641_simple_0.phy[0]', 'mpc8641_simple_0.cell')
                   ('mpc8641_simple_1.phy[0]', 'mpc8641_simple_1.cell')
simics> c

```

It should now be possible to ping between the two simulated machines. Enter the following commands in the first machine.

```

~ # ping -c 3 10.10.0.41
PING 10.10.0.41 (10.10.0.41): 56 data bytes
64 bytes from 10.10.0.41: seq=0 ttl=64 time=20.226 ms
64 bytes from 10.10.0.41: seq=1 ttl=64 time=20.217 ms
64 bytes from 10.10.0.41: seq=2 ttl=64 time=20.216 ms

--- 10.10.0.41 ping statistics ---
3 packets transmitted, 3 packets received, 0% packet loss
round-trip min/avg/max = 20.216/20.219/20.226 ms
~ #

```

Network simulation is covered in more details in [chapter 3](#).

2.2 Connect to a Real Network

A network simulation can also be connected to a real network. By doing this, simulated computers and real computers will be able to communicate with each other.

Note: Timing problems when using a real network connection, such as TCP timeouts, indicates that the simulation is running too fast: it might be idle most of the time. In this case, slow it down by using the **enable-real-time-mode** command.

Before following the steps in this example, start a new `firststeps.simics`, but do not boot it yet. Connecting a simulated machine to a real network can be as simple as typing a single command:

```

simics> connect-real-network
Created instantiated 'std-service-node' component 'default_service_node0'

```

2.2. Connect to a Real Network

```
NAPT enabled with gateway 10.10.0.1/24 on link default_eth_switch0.link.
NAPT enabled with gateway fe80::2220:20ff:fe20:2000/64 on link default_eth_switch0.link.
Host TCP port 4021 -> 10.10.0.40:21
Host TCP port 4022 -> 10.10.0.40:22
Host TCP port 4023 -> 10.10.0.40:23
Host TCP port 4080 -> 10.10.0.40:80
Real DNS enabled at 10.10.0.1/24 on link default_eth_switch0.link.
Real DNS enabled at fe80::2220:20ff:fe20:2000/64 on link default_eth_switch0.link.
```

Note: Note that **connect-real-network** found out which IP address to work with by checking configuration variables set by `firststeps.simics`. You can also specify this address as parameter to the command

The command above creates a new Ethernet switch, and connects it to all simulated network devices as well as to the real network. It also enables NAPT, *Network Address Port Translation*. Finally, it forwards ports 4021, 4023 and 4080 on the simulation host to the simulated machine's telnet, FTP and HTTP ports, respectively.

To make everything work in practice, the simulated system must be configured to know about the outside world. Boot the machine, then issue the following commands to set up a gateway and a domain name server:

```
~ # route add default gw 10.10.0.1
~ # echo nameserver 10.10.0.1 > /etc/resolv.conf
~ #
```

The simulated system is now directing all accesses outside the local network to the gateway at 10.10.0.1, which is the *service node* that **connect-real-network** created for us. The service node will act as a router to connect to other networks, including the real ones.

Now, if the host computer used for the simulation is connected to Internet, it is possible to telnet to a real computer. In this example we use `gnu.org`.

```
~ # telnet gnu.org 80
GET /
<?xml version="1.0" encoding="ISO-8859-1"?>
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"
    "http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">
<html xmlns="http://www.w3.org/1999/xhtml" lang="en" xml:lang="en">
<head>
<title>Server error!</title>
[...]
```

Connection closed by foreign host.

```
~ #
```


2.2. Connect to a Real Network

Since **connect-real-network** has set up forwarding for the telnet, FTP and HTTP ports of the simulated machine, it is possible to telnet into the simulated machine. On your host, try the following:

```
host:~$ telnet localhost 4023
Trying 127.0.0.1...
Connected to localhost.
Escape character is '^]'.

(none) login: root
~ # ls /
bin          etc          host          lib           lost+found    root          sys
dev          home         init          linuxrc       proc          sbin          usr
~ # exit
Connection closed by foreign host.
```

You can also access the minimal web server running on the simulated computer. Point your browser to the following address: `http://localhost:4080`. You should get something similar to figure 2.1.

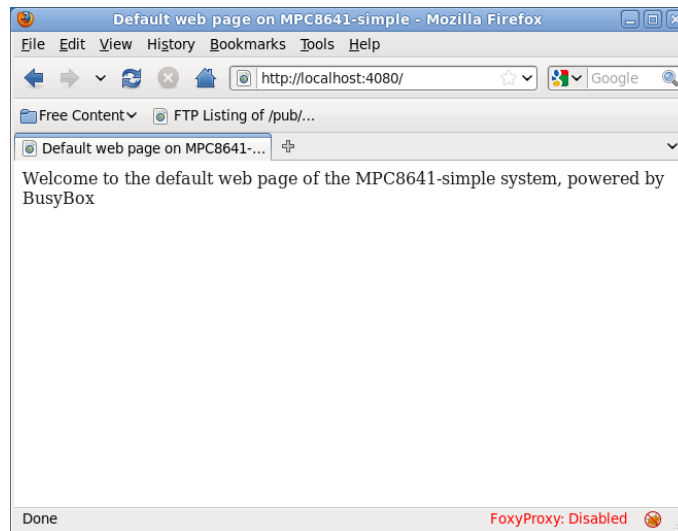


Figure 2.1: A very simple page

Chapter 4 is the reference for everything related to real network connections.

Chapter 3

Network Simulation

This chapter describes how Simics models Ethernet networking.

3.1 Ethernet Links

Connecting simulated machines over a simulated Ethernet connection is done by creating a *link* using one of the components in the **eth-links** module: **ethernet-cable**, **ethernet-hub** or **ethernet_switch**. Any of these components can connect to the Ethernet devices of the simulated machines. They model Ethernet at the frame level, in that they perform delivery of complete frames sent from one device to another, and do not model collisions or lower-level signaling details. Each of the three ethernet links have their own particularities:

ethernet_cable

An Ethernet cable can only be connected to two devices, since, as a real ethernet cable, it only has two ends. Apart from transporting ethernet frames back and forth, an ethernet cable models whether the signal is up or down.

ethernet_hub

An **ethernet_hub** is the simplest ethernet link model: all frames are broadcast to all devices connected to the hub. It works like a traditional hub or coaxial cable.

ethernet_switch

An **ethernet_switch** also allows the connection of multiple devices, but functions rather like a switch: as it learns what MAC addresses the different devices have, it stops broadcasting and transports the frame more selectively. This can substantially improve simulation performance when several devices are talking on the link, and even more so when the simulation is distributed.

As all three models work more or less at the frame level; they can be used for any type of Ethernet connection, without limitations due to speed or technology. Traffic sent over the link can be anything, including TCP/IP or any other protocol stack that works on top of Ethernet. The link does not need to be configured for the intended use.

The three components are provided in the **eth-links** module. They can be manipulated in the same way as any other component: after importing **eth-links**, the **create-ethernet-cable/hub/switch** commands are made available. They create non-instantiated components:

```
simics> load-module eth-links
simics> create-ethernet-hub
Created non-instantiated 'ethernet-hub' component 'ethernet_hub0'.
```

The name will be automatically chosen by Simics unless it is provided as an argument to the command:

```
simics> create-ethernet-hub my_ethernet_hub
Created non-instantiated 'ethernet-hub' component 'my_ethernet_hub'.
```

When a link component has been created, it is possible to connect the network devices by running the **connect** command for each Ethernet port. We will take here the *MPC8641D* board from the Firststeps package as an example. Start the simulation with the `firststeps.simics` start script:

```
simics> load-module eth-links
simics> create-ethernet-hub my_ethernet_hub
Created non-instantiated 'ethernet_hub' component 'my_ethernet_hub'
simics> connect mpc8641d_simple.eth[0] my_ethernet_hub.device0
simics>
```

When all of the components of the system are configured and connected together, instantiate the link components:

```
simics> instantiate-components
```

It is also possible to create instantiated links with the **new-ethernet-cable/hub/switch** commands, but note that this only works if there is already at least one instantiated top-level component in the simulation.

3.2 Ethernet VLAN Switch

The **ethernet_switch** described in the previous section complains if it receives frames tagged with VLAN information. For VLAN the **eth-links** provides a IEEE 802.1Q VLAN aware switch: **ethernet_vlan_switch**.

Note: This section assumes you know what VLAN is. It only shows how to configure and use the `ethernet_vlan_switch`.

The VLAN switch has the same functionality as the basic Ethernet switch, but the interface is different to handle the added complexity of VLAN. You need to add VLAN groups to the switch before you can connect Ethernet devices to it. Each VLAN group has its own set of connectors.

The VLAN switch is created in the same way as the other Ethernet link components:

```
simics> load-module eth-links
simics> create-ethernet-vlan-switch evs
Created non-instantiated 'ethernet_vlan_switch' component 'evs'
```

Once the link has been created you need to add VLAN groups to it:

```
simics> evs.add-vlan 1
```

Each VLAN group adds two kinds of connectors: `vlan_X_devY` and `vlan_X_trunk_devY`, where `X` is the VLAN group id and `Y` is the an identifier for the connector. When a connector of one kind is used the switch automatically creates a new empty connector of the same kind with a new identifier.

Use the `vlan_X_devY` connectors to connect devices which do not expect or add VLAN tags. For example, to connect the first network port of the MPC8641 machine created by the `firststeps` script:

```
simics> connect mpc8641d_simple.eth[0] evs.vlan_1_dev0
```

Frames sent from the switch on these connections will not have any VLAN tags and the switch will generate a warning if it receives a tagged frame on such a connection.

Each VLAN group also has trunk connectors. On these connections the switch will parse frames sent to the switch to find which VLAN group they should be sent to and the switch will ensure that outgoing frames are tagged with the correct VLAN group.

To avoid having to keep track of the connector identifiers you can use **get-free-connector** to get an available non-trunk connector for a VLAN group:

```
simics> evs.get-free-connector 1
"evs.vlan_1_dev1"
simics> connect mpc8641d_simple.eth[1] (evs.get-free-connector 1)
simics> evs.get-free-connector 1
"evs.vlan_1_dev2"
```

and **get-free-trunk-connector** to get an available trunk connector for a VLAN group:

```
simics> evs.get-free-trunk-connector 1
"evs.vlan_1_trunk_dev0"
```

3.3 Link Object Timing

All frames that are sent over a link are delivered to the receiving devices after a small delay. The delay is the same for every frame, and is called the *latency* of the link. Each link object has a *goal_latency* configuration parameter that controls the ideal latency the link wants to obtain.

Link objects are most often used to communicate between network devices using separate clocks. Due to way how Simics handles simulated time, different clocks are not always completely synchronized. In order to avoid indeterministic simulation, the link latency must be high enough that any data sent over the link will never reach the recipient at a point in time it has already passed. This imposes a lower boundary on the latency, called the *minimum latency* of the link. The value of the minimum latency depends on the simulation setup, in particular whether the simulation is multithreaded or distributed across several Simics processes. See the *Accelerator User's Guide* for more information about links and latencies.

The latency of a link can be specified in the **create-ethernet-cable/hub/switch** command as a time in seconds. If the latency of a link is set too low, it will be automatically adjusted to the lowest value allowed by the setup when the component is instantiated. For example, when creating two instances of the **MPC8641D** machine and connecting them to an Ethernet link with a too low latency, Simics will adjust the latency automatically. Let us look at an example with the `firststeps-multi.simics` script. It uses the **new-** command to create instantiated components instead of the **create-** variant:

```
simics> load-module eth-links
simics> new-ethernet-hub goal_latency = 0.000001
Created instantiated 'ethernet_hub' component 'ethernet_hub0'
simics> connect mpc8641_simple_0.eth[0] ethernet_hub0.device0
simics> connect mpc8641_simple_1.eth[0] ethernet_hub0.device1
```

The *effective latency* of a link can be displayed by the `<ethernet-link>.status` command. The actual Ethernet links are not created until the components are instantiated. To get access to the Ethernet link implementation, just request the **link** object in the component:

```
simics> ethernet_hub0.link.status
Status of ethernet_hub0.link [class eth-hub-link]
=====

Effective latency : 10.0 ms

Connected devices : ('mpc8641_simple_0.phy[0]', 'mpc8641_simple_0.cell')
                  ('mpc8641_simple_1.phy[0]', 'mpc8641_simple_1.cell')
```

Note the higher effective latency the link obtained, despite a goal latency of 1us. The default latency can be adjusted using the **set-min-latency** command.

3.4 IP Services

It is often useful to let the simulated machines use services available on the simulated network, especially at boot time. Normally these services runs on servers connected to the network. To avoid having to set up simulated servers just to provide them, Simics implements a *service node* instead.

The **std-service-node** component class, available after importing the **std-components** module, provides a virtual network node that acts as a server for a number of TCP/IP-based protocols, and as an IP router between simulated networks. It handles both IPv4 and IPv6 protocol versions. The supported services are:

- IP based Routing (v4 and v6)
- RARP (v4)
- DHCP/BOOTP (v4) and DHCPv6 (v6)
- DNS (v4 and v6)
- FTP (v4)
- TFTP (v4)
- Real network connections (see chapter 4)

There can be any number of **std-service-node** components, and each one can be connected to any number of Ethernet links. In most configurations, however, there will be a single service node. A service node can be created using the **create/new-std-service-node** commands:

```
simics> new-std-service-node sn0
Created instantiated 'std-service-node' component 'sn0'.
```

This service node can then be connected to an Ethernet link component. Note that the service-node IP address on the link must be specified:

```
simics> sn0.connect-to-link ethernet_hub0 10.10.0.1
sn0.connect-to-link ethernet_hub0 10.10.0.1
Adding host info for IP 10.10.0.1: simics0.network.sim
MAC: 20:20:20:20:20:00
simics>
```

The rest of the configuration can be done *when the service node has been instantiated*, using the available commands, such as:

```
simics> sn0.list-host-info
IP          name.domain          MAC
-----
10.10.0.1   simics0.network.sim    20:20:20:20:20:00
simics>
```

3.4.1 IP Based Routing

A **service-node** can provide IP based routing between Ethernet links, allowing machines attached to different networks to communicate with each other.

Note: To use the routing mechanisms, simulated machines must use the IP address of the service node as a *gateway* for IP based traffic. Configuring a gateway requires system administration skills, and the exact procedure depends on the target operating system.

Each connection of the service-node to an Ethernet link implies a default route to that link. For example, connecting a service node with the address 192.168.0.1/24 to **link1** implies that all packets matching this network and mask combination will be routed to **link1** automatically. This often solves the most common routing needs.

In addition, the service node contains an internal IP routing table that is used for packet routing between connected links. The routing table can be viewed using the **(service-node).route** command:

```
simics> sn0.route
Destination  Netmask  Gateway  Link
-----
10.10.0.0    24              link0
```

The output is quite similar to `route` command available on many systems. The *destination* and *netmask* fields specify a target that can be either a network (i.e., a range of addresses) or a single host (with netmask 255.255.255.255). For packets with this target as their destination, the *link* field specifies the Ethernet link the packet should be sent to.

New entries can be added to the routing table with the **(service-node).route-add** command. If there is a service node called **sn0** connected to two links called **link0** and **link1**, it would for example possible to set up routes like this:

```
simics> sn0.route-add 192.168.0.0 255.255.0.0 link = link0
simics> sn0.route-add 192.168.1.0/26 link = link1
simics> sn0.route-add 10.10.0.0 255.255.0.0 192.168.0.1 link0
simics> sn0.route-add 0.0.0.0 255.255.255.255 192.168.1.1 link1
simics> sn0.route
Destination  Netmask  Gateway      Link
-----
192.168.0.0  16              link0
192.168.1.0  26              link1
10.10.0.0    16      192.168.0.1  link0
default      16      192.168.1.1  link1
```

The destination address and the netmask identify the target, and should be given as strings in dotted decimal form. If the target is a single host, the netmask should be given as "255.255.255.255".

3.4.2 DHCP and BOOTP

A service node can act as a *Dynamic Host Configuration Protocol* (DHCP) or *Bootstrap Protocol* (BOOTP) server, responding to requests from clients that can read their network configuration from such a server. The DHCP protocol is an extension of the BOOTP protocol, and for many uses the feature set is more or less the same. The Simics implementation uses the same configuration for both services.

The service node has a table that maps MAC addresses to IP addresses and domain name. This is used to answer DHCP or BOOTP requests. The `<service-node>.add-host` command can add entries to this table:

```
simics> sn0.add-host 10.10.0.1 node1 mac="10:10:10:10:10:01"
Adding host info for IP 10.10.0.1: node1.network.sim
MAC:10:10:10:10:10:01
```

The `<service-node>.list-host-info` command prints the current contents of the table:

```
simics> sn0.list-host-info
IP          name.domain          MAC
-----
10.10.0.0   simics0.network.sim  20:20:20:20:20:00
10.10.0.1   node1.network.sim    10:10:10:10:10:01
```

The `<service-node>.dhcp-add-pool` command adds dynamic DHCP leases, from which new clients will be automatically assigned an address on request. When an entry from the pool is given out, the new mapping is stored in the internal host info table, including a generated name that can be found through DNS queries. If a DHCP client's MAC address matches an entry in the table, it is assigned the corresponding IP address. If there is no matching MAC address, the dynamic address pools will be searched for an available IP address.

The DHCP implementation in **service-node** is simple, and might not work with all DHCP clients.

3.4.3 DNS

The service node includes the functionality of a simple *Domain Name Server* (DNS), that a simulated client can use to translate a host/domain name into an IP address and vice versa. The DNS service is based on the same host table as the DHCP service, and only answers requests for A and PTR records.

For entries in the table that will only be used to answer DNS requests, and not for DHCP, the MAC address can be left out. The `<service-node>.add-host` command can be used to add table entries, and the `<service-node>.list-host-info` command prints the current table. By default, all host entries will use the `network.sim` domain.

```
simics> sn0.add-host 10.10.0.1 donut
Adding host info for IP 10.10.0.1: donut.network.sim
```



```
simics> sn0.add-host 10.11.0.1 foo other.domain
Adding host info for IP 10.11.0.1: foo.other.domain
simics> sn0.list-host-info
```

IP	name.domain	MAC
10.10.0.0	simics0.network.sim	20:20:20:20:20:00
10.10.0.1	donut.network.sim	10:10:10:10:10:01
10.11.0.1	foo.other.domain	

For dynamic DHCP addresses, a DNS name will be added for the new IP number, so that any of these addresses can be found by the DNS service. When connected to a real network, the DNS service can do external lookups for names it does not recognize.

3.4.4 TFTP

The service node also supports the *Trivial File Transfer Protocol* (TFTP, see RFC 1350) which allows to transfer files between the host system (running the simulation) and a simulated (target) client. TFTP is often used during network booting, together with the BOOTP facilities, to load OS kernels and images, and it can also be used interactively with the `tftp` command found on many systems.

Files to be transferred from the host system to the simulated client should be placed in a directory in the Simics path. This is the standard path used by image objects: **list-directories** will print its current value, while **add-directory** will add a directory to the path list. The current working directory is also automatically included.

Files transferred from the simulated client to the host will end up in the current working directory. When running Simics in graphical mode, this will be the workspace directory; from the command line or without a workspace, it will be the directory Simics was started from.

Note: TFTP is based on UDP, and each packet is acknowledged individually before the transfer is allowed to continue. Depending on the latency of the link, the transfer of large files can be slow. In that case, ensuring that the link uses a lower latency will increase performance.

A short example, assuming the target machine is running Linux with TFTP installed, has already booted, and has a working service node properly connected. We bring up the target's network interface:

```
joe@computer: ~# ifconfig eth0 10.10.0.10 up
```

And we transfer the file `myfile.txt` from the host machine:

```
joe@computer: ~# tftp -l myfile.txt -r myfile.txt -g 10.10.0.1
```

3.5 Observing Network Traffic

Simics proposes several ways of listening to the network traffic on one, or all Ethernet links. Simics includes ready-to-use traffic dumping capabilities, using external network monitoring tools. It also provides an Ethernet probe and Ethernet snoop capabilities, giving access to a programmatic interface for listening to the network traffic.

3.5.1 Traffic Monitoring Software

Simics provides several commands to dump the traffic on one or several Ethernet links. These commands use existing file formats and network monitoring tools to present the results:

pcap-dump

The `<ethernet_link>.pcap-dump` command dumps the traffic of a specific Ethernet link to a file that can be read by `tcpdump` or any compatible program like `Wireshark` (www.wireshark.org). A global **pcap-dump** also exists to dump the traffic of all Ethernet links of the simulation in the same file.

tcpdump

The **tcpdump** command works as **pcap-dump**, but the traffic is redirected to a `tcpdump` instance running in a separate window (for Unix only).

wireshark

The **wireshark** command works as **tcpdump**, but redirects the traffic to an instance of Wireshark instead.

Note that there can only be one traffic dumping tool active on the link. Simics will automatically stop the current traffic dump and start a new one as necessary. The **pcap-dump-stop**, **tcpdump-stop**, and **wireshark-stop** commands can be used to stop the traffic dumping.

Note: When using **std-ethernet-link** components rather than the new-style components, traffic dumping is still supported but the traffic of new links and old links can not be mixed. For example, if both types of links exist, two `tcpdump` instances will be started, one for the new links, and one for the old ones.

These commands can also be activated for a specific device on the link by associating them with an existing Ethernet probe. This is described in the *Ethernet Probe* section below.

3.5.2 Ethernet Probe

Note: This section only applies to new-style links, such as **ethernet-switch**, **ethernet-hub**, or **ethernet-cable**, but not **std-ethernet-link**.

The Ethernet probe provides a way to listen to traffic at a particular endpoint of the link, that is, the probe will receive both incoming and outgoing traffic for a particular device.

A probe is inserted using the **insert-ethernet-probe** command with appropriate arguments. Using *Firststeps* as an example machine, the commands would be:

```

simics> load-module eth-links
simics> new-ethernet-switch link0
Created instantiated 'ethernet_switch' component 'link0'
simics> connect link0.device0 mpc8641d_simple.eth[0]

# insert a probe between the PHY of the eth[0] device above and the link
simics> load-module eth-probe
simics> insert-ethernet-probe device = mpc8641d_simple.phy[0]
Created probe 'probe0'
simics> probe0.info
Information about probe0 [class eth-probe]
=====

Connections:
  Port A : mpc8641d_simple.phy[0]
  Port B : link0.link

```

At this point, the probe is ready to use. You can issue a `<eth-probe>.pcap-dump` or similar command to connect an external network monitoring tool at the probe level. The traffic will be dumped as seen from the `mpc8641d_simple.phy[0]` device.

You can also register your own callback to listen to the traffic going-on in the probe, using the `ethernet_probe` interface provided by the probe object:

```

typedef enum {
    Eth_Probe_Port_A = 0,
    Eth_Probe_Port_B = 1
} eth_probe_side_t;

typedef void (*ethernet_probe_snoop_t)(lang_void *user_data,
                                       conf_object_t *probe,
                                       eth_probe_side_t to_side,
                                       const frags_t *frame,
                                       eth_frame_crc_status_t crc_status);

SIM_INTERFACE(ethernet_probe) {
    void (*attach_snooper)(conf_object_t *NOTNULL probe,
                          ethernet_probe_snoop_t snoop_fun,
                          lang_void *user_data);
    void (*attach_probe)(conf_object_t *NOTNULL probe,
                        ethernet_probe_snoop_t snoop_fun,
                        lang_void *user_data);
    void (*detach)(conf_object_t *NOTNULL probe);
}

```

```

void (*send_frame)(conf_object_t *NOTNULL probe,
                   eth_probe_side_t to_side,
                   const frags_t *frame,
                   eth_frame_crc_status_t crc_status);
};

#define ETHERNET_PROBE_INTERFACE "ethernet_probe"

```

A complete description of this interface is provided in the *Hindsight Reference Manual*. What we are interested in at this point is to register a *snooper* callback that will only listen to traffic:

```

# a callback that does nothing but print a warning
simics> def callback(user_data, probe, to_side, packet, crc_status):
    if to_side == Eth_Probe_Port_A:
        print 'packet going to device'
    else:
        print 'packet going to network'

.....
simics> @conf.probe0.iface.ethernet_probe.attach_snooper(callback, None)
simics> c
[mpc8641d_simple.soc.mcm info] Enabling core 1
packet going to network
packet going to network
packet going to network
[...]
```

The probe can also drop, modify or inject packets. This is described in the *Injecting Network Traffic* section below.

3.5.3 Traffic Snooping

Note: This section only applies to new-style links, such as **ethernet-switch**, **ethernet-hub**, or **ethernet-cable**, but not **std-ethernet-link**.

Ethernet links provides a special interface to listen to all traffic on the link via a function callback. This makes it possible to write simple traffic dumping scripts with customized output.

```

typedef void (*ethernet_link_snoop_t)(lang_void *user_data,
                                       conf_object_t *clock,
                                       const frags_t *packet,
                                       eth_frame_crc_status_t crc_status);

SIM_INTERFACE(ethernet_snoop) {

```

```

conf_object_t *(*attach)(conf_object_t *NOTNULL link,
                          conf_object_t *clock,
                          ethernet_link_snoop_t snoop_fun,
                          lang_void *user_data);
};
#define ETHERNET_SNOOP_INTERFACE "ethernet_snoop"

```

This interface is implemented by ethernet link objects. It is used to attach snoop functions to the link. The snoop function will receive all traffic going over the link.

This interface should only be used for inspection, and never as part of the actual simulation. The snoop functions must not affect the simulation in any way.

The *clock* parameter tells the link on which clock to post the events that call the snoop function. The snoop function will be called at the delivery time of the network packet, which means that it will be called at the same time as any ethernet devices attached to the same clock that receives packets from the same link.

Snooped frames with a matching CRC will contain the correct frame check sequence.

The *user_data* parameter is passed to the snoop function every time it is called.

Using the *Firststeps* as an example, a simple script callback could be written as follow:

```

simics> load-module eth-links
simics> new-ethernet-switch link0
Created instantiated 'ethernet_switch' component 'link0'
simics> connect link0.device0 mpc8641d_simple.eth[0]

# a callback that does nothing but print a warning
simics> @def callback(user_data, clock, packet, crc_status):
simics>     print "packet received in snooper"

.....

# callback registration on link0, using the first CPU as clock object
simics> @ep = conf.link0.link.iface.ethernet_snoop.attach(
conf.mpc8641d_simple.soc.cpu[0], callback, None)
simics> c
[mpc8641d_simple.soc.mcm info] Enabling core 1
packet received in snooper
packet received in snooper
packet received in snooper
[...]
```

The endpoint object returned by the *attach()* function can be destroyed at any time using *SIM_delete_object()*, ending the capture. Snooper endpoints are used by the external monitoring tools system described in the previous section to feed to the tools the packets passing on the links.

3.6 Injecting and Modifying Network Traffic

3.6.1 From a Network Dump

Note: This section only applies to new-style links, such as **ethernet-switch**, **ethernet-hub**, or **ethernet-cable**, but not **std-ethernet-link**.

Simics comes with a **eth-injector** class that takes a *pcap* format file and injects the packets described in the file into the simulated network.

Each packet in a *pcap* file has a time stamp that usually is the absolute time when the packet was recorded. The **eth-injector** starts injecting the first packet of the *pcap* file directly after the **start** command has been run. The consecutive packets are injected after an amount of virtual time that is equal to the difference in time stamp between that packet and the first packet of the *pcap* file. If the packet cannot be injected because of bandwidth limitations, it is ignored. Incoming packets are ignored as well.

It is common that the CRC of Ethernet frames are not recorded in *pcap* files. In Simics, the whole Ethernet frame has to be present for a correct simulation result. Using the *-no-crc* option when running the **<eth-injector>.start** command tells the injector that the *pcap* file contains no CRC. The injector then adds a CRC to each frame that Simics will handle as if it was correct. By not using the *-no-crc* option the frames in the *pcap* file are injected as they were recorded, without any modification.

The **eth-injector** can be connected to an Ethernet link like any other Ethernet device. It can also be connected directly to another Ethernet device without the need to have a link between the device and the **eth-injector**.

The following example creates a new **eth-injector**, connects it to an already existing **ethernet_switch** of name **ethernet_switch0** and starts packet playback from a file named *test.pcap*:

```
simics> load-module eth_injector_comp
simics> new-eth-injector-comp name = inj0
Created instantiated 'eth_injector_comp' component 'inj0'
simics> ethernet_switch0.connect component = inj0
simics> inj0.injector.start file = test.pcap
```

3.6.2 From an Ethernet Probe

Note: This section only applies to new-style links, such as **ethernet-switch**, **ethernet-hub**, or **ethernet-cable**, but not **std-ethernet-link**.

The section *Observing Network Traffic* explained how to use an Ethernet probe to listen to the incoming and outgoing traffic of a device connected on an Ethernet link. A probe can also be used to modify this traffic by dropping and changing existing packets, or injecting new ones.

3.6. Injecting and Modifying Network Traffic

```
typedef enum {
    Eth_Probe_Port_A = 0,
    Eth_Probe_Port_B = 1
} eth_probe_side_t;

typedef void (*ethernet_probe_snoop_t)(lang_void *user_data,
                                       conf_object_t *probe,
                                       eth_probe_side_t to_side,
                                       const frags_t *frame,
                                       eth_frame_crc_status_t crc_status);

SIM_INTERFACE(ethernet_probe) {
    void (*attach_snooper)(conf_object_t *NOTNULL probe,
                          ethernet_probe_snoop_t snoop_fun,
                          lang_void *user_data);
    void (*attach_probe)(conf_object_t *NOTNULL probe,
                        ethernet_probe_snoop_t snoop_fun,
                        lang_void *user_data);
    void (*detach)(conf_object_t *NOTNULL probe);
    void (*send_frame)(conf_object_t *NOTNULL probe,
                      eth_probe_side_t to_side,
                      const frags_t *frame,
                      eth_frame_crc_status_t crc_status);
};

#define ETHERNET_PROBE_INTERFACE "ethernet_probe"
```

A complete description of this interface is provided in the *Hindsight Reference Manual*. What we are interested in now is to register a *probe* callback that will be allowed to modify the traffic:

```
# a callback that drops all outgoing packets
simics> @def callback(user_data, probe, to_side, packet, crc_status):
    if to_side == Eth_Probe_Port_A:
        print 'dropping incoming packet'
    else:
        print 'forwarding outgoing packet'
        probe.iface.ethernet_probe.send_frame(to_side, packet, crc_status)

.....
simics> @conf.probe0.iface.ethernet_probe.attach_probe(callback, None)
simics> c
```

3.6. Injecting and Modifying Network Traffic

```
[mpc8641d_simple.soc.mcm info] Enabling core 1  
forwarding outgoing packet  
forwarding outgoing packet  
forwarding outgoing packet  
[...]
```

When a callback is registered as a *probe*, it takes responsibility for forwarding packets it receives. It is also allowed to drop them, modify them, modify their CRC status or inject new packets. For example, it could duplicate outgoing packets, inject errors, etc. Note that if your callback modifies the simulation in this manner, you may need to create an object representing the changes' state engine to make sure the simulation can be checkpointed and stays deterministic.

Chapter 4

Connecting to a Real Network

Connecting a simulator to a real network opens many new possibilities. For example, it makes it is easy to download files to the simulated machines using FTP, to access the simulated machines remotely through telnet, or to test software on simulated machines against real machines.

Simics provides four kinds of real-network connections. This chapter will go through all of them in details, with numerous examples. Section 4.1 starts by explaining the preparatory steps for all the examples presented afterward. Section 4.2 goes through all four kinds of real-network connections, how they are set up and what they provide.

Some real-network connections require specific configuration on the host system: these are covered in section 4.3. Section 4.4 describes how to select the correct host interface when running on a host with multiple Ethernet interfaces.

Finally, section 4.5 describes how to tune real network connections to improve their performance, and section 4.6 contains a troubleshooting guide to help diagnose problems.

Note: Connecting a simulated network to a real network requires some knowledge of network administration issues.

4.1 Preparing for the Examples

All the examples in section 4.2 use the simulated *MPC8641D* machine from the First Steps tutorial. Start the `firststeps.simics` configuration and boot the machine. Then, to save time when trying several examples, save a checkpoint. This checkpoint can be used instead of launching and booting *firststeps* for each example. When idling, *firststeps* runs faster than real time, which means that it can time out faster than expected, unless real time mode is turned on with **enable-real-time-mode**.

Some of the examples involve using *telnet* to connect to the simulated machine. On both Unix and Windows, a telnet binary is provided and can be used directly respectively in a terminal or in a Command-Line prompt. The output on Windows may be slightly different from the output given in the examples.

In the examples, the simulated machine sometimes needs to be reconfigured. Since it is running Linux, the appropriate Linux configuration commands are indicated. Other operating systems should be configured similarly, but the commands may of course differ.

Finally, in all examples, the host machine where Simics is running has the IP address 10.0.0.129 and the real host that communicates with the simulated network has the IP address 10.0.0.240. These addresses should be replaced as necessary. *Firststeps* uses its default IP address 10.10.0.40.

The examples assume that there is a host on the real network that accepts telnet connections. Check that it is possible to telnet from the simulation host to the other real host. Just run `telnet ip` at a command-line, where *ip* is the IP address of the other real host. If that does not work, the simulated machine will not be able to connect to the real host either.

If there is no host that accepts telnet connections on the network, the connection can be tested with a web server on port 80 instead, by entering `GET / HTTP/1.0` and a blank line. This should return the HTML content of the start page of the server. Here `www.google.com` is used:

```
~ # telnet www.google.com 80
Trying 64.233.161.104...
Connected to 64.233.161.104.
Escape character is '^]'.
GET / HTTP/1.0

HTTP/1.0 302 Found
Location: http://www.google.se/cxfer?c=PREF%3D:TM%3D1118841789:S%3DumC
Vbug84n5uBWAo&prev=/
Set-Cookie: PREF=ID=a5e237e2402bdcac:CR=1:TM=1118841789:LM=1118841789:
S=HQ3jOc8_lpeVGj98; expires=Sun, 17-Jan-2038 19:14:07 GMT; path=/;
domain=.google.com
Content-Type: text/html
Server: GWS/2.1
Content-Length: 214
Date: Wed, 15 Jun 2005 13:23:09 GMT
Connection: Keep-Alive

<HTML><HEAD><TITLE>302 Moved</TITLE></HEAD><BODY>
<H1>302 Moved</H1>
The document has moved
<A HREF="http://www.google.se/">here</A>.
</BODY></HTML>
Connection closed by foreign host.
~ #
```

Make sure that the telnet or web server is on the same IP subnet as the simulation host, since it may not be possible to access other subnets, depending on what real-network connection is in use.

4.2 Connection Types

There are four kinds of connections between simulated networks and real networks in Simics. The next paragraphs describe how they work, and their advantages and drawbacks.

All connection types except port forwarding require low-level access to the simulation host's Ethernet interfaces, and therefore require administrative privileges to set up. However, administrative privileges are, in most cases, not needed once the low-level access has been set up. See section 4.3 for details.

Port forwarding

Port forwarding is the easiest connection type to set up for simple usage. It does not require administrative privileges nor any configuration on the simulation host or on the other hosts. Port forwarding is part of the Simics Hindsight product and does not require the Simics Ethernet Networking product.

However, port forwarding is limited to TCP and UDP traffic. Other traffic, for example, ping packets that use the ICMP protocol, will not pass through the connection. Since port forwarding uses ports on the simulation host it is not possible to use incoming ports that are already used by the simulation host, or ports below 1024 without administrative privileges.

Each incoming TCP port, and each incoming or outgoing UDP port require a separate forwarding rule. Therefore, for an application that uses many ports, or random ports, configuration can become cumbersome or nearly impossible without complex communication. Outgoing TCP connections on many or random ports can be handled by NAT, so that is not a problem.

Port forwarding allows communication between the simulated machines, the simulation host and other hosts on the real network.

Ethernet bridging connection

With an Ethernet bridging connection, the simulated machines appears to be directly connected to the real network. The connection allows any kind of Ethernet traffic between the simulated and real networks. Usually IP addresses from the IP subnet of the real network are used by the simulated machines, in which case nothing needs to be configured on the real hosts on the real network. However, the simulation host can not be accessed from the simulated machines using an Ethernet bridging connection.

To use Ethernet bridging, the simulation host needs to be set up for TAP access as described in section 4.3.

This connection type is provided in the Simics Ethernet Networking product.

Host connection

With a host connection, the simulation host is connected to a simulated network, allowing any kind of Ethernet traffic between the simulation host and the simulated machines.

Host connections also supports *IP forwarding*. When using IP forwarding, the operating system of the host routes IP traffic between the real and simulated networks. As above, routes should be configured between the simulated and real networks to make it work.

To use host connections, the host needs to be set up for TAP access as described in section 4.3.1.

This connection type is provided in the Simics Ethernet Networking product.

Figure 4.1: Comparison of real-network connections

	Port F	Bridge T	Bridge NT	Host
Need Ethernet Network product	no	yes	yes	yes
Need admin rights for config.	no	yes	yes	yes
Need admin rights to run	no	no	yes (1)	no
Need real IP available	no	yes	yes	no
Support UDP/TCP	yes	yes	yes	yes
Restrict TCP/UDP ports	yes	no	no	no
Support all IPv4	no	yes	yes	yes
Support all Ethernet	no	no	yes	yes (2)
Talk to simulation host	yes	no	no(5)	yes
Need config on real hosts	no	no	no	yes (3)
Need dedicated network	no	no	yes (4)	no

- **PF:** port-forwarding
- **Bridge T:** ethernet bridging with mac translation
- **Bridge NT:** ethernet bridging without mac translation
- **Host:** host connection

1. The ethernet interface needs to be put in promiscuous mode.
2. Only with the simulation host.
3. Only on simulation host.
4. Needed if there is too much unrelated traffic on real network.
5. Possible using TAP access, but not recommended.

The table 4.1 recapitulates the advantages and drawbacks of each type of connection. Basically, for simple TCP services like FTP, HTTP or telnet, port forwarding is the way to go. If port forwarding does not suffice and if there are available IP addresses on the IP subnet of the real network, or for network protocols other than IPv4, Ethernet bridging is another possibility. Finally, if access to the simulated machines from the simulation host is required, but port forwarding is not sufficient, host connection might be the solution.

All commands that create a connection to the real network start with the prefix **connect-real-network-**, with different suffixes depending on the connection type. They come in two variants.

For each connection type there is a global command that assumes that there is at most one Ethernet link object. If there is no Ethernet link object, a default **ethernet_switch** is created. All Ethernet interfaces of all simulated machines in the Simics process are then automatically connected to the new Ethernet switch, and the Ethernet switch is connected to the real network. This is an easy way to connect all simulated machines in the Simics process to the real network with a single command. For example, to connect all simulated machines to the real network using an Ethernet bridging connection, just type in the global command **connect-real-network-bridge**.

For a more complex simulated network setup, not all simulated Ethernet interfaces will be connected to the same network. In that case, create first the simulated network setup, and then connect specific Ethernet links to the real network. For each connection type, there is a command with the same name as the global command that can be run on a specific Ethernet link object to connect it to the real network. For example, with an Ethernet link object named **ethernet_hub0**, use the command **ethernet_hub0.connect-real-network-bridge** to create an Ethernet bridging connection between that particular link and the real network.

The commands related to port forwarding are an exception to this rule. They do not come in variants that can be run on Ethernet links objects, but instead have an *ethernet-link* argument that can be used to specify a link.

4.2.1 Port Forwarding

Port forwarding forwards traffic on TCP and UDP ports between the simulated network and the real network. It also allows forwarding DNS queries from the simulated network to the real network. Port forwarding can be used with any kind of IP network on the host, it is not limited to Ethernet networks.

Port forwarding is probably the easiest way to access the real network for simple TCP or UDP connectivity, for example, telnet or FTP usage. Port forwarding is easy to set up. Simics does not need administrative privileges to run port forwarding, and neither the simulation host nor any other host needs to be configured in any way.

Port forwarding is managed by a service node connected to an Ethernet link. It is the service node that listens for traffic on both the real and simulated networks and forwards it to the other side. All port forwarding commands except **connect-real-network** therefore take as argument an Ethernet link with a connected service node.

There are really four distinct parts to Simics's port forwarding solution: forwarding of specific ports from the real network to the simulated network, forwarding of specific ports from the simulated network to the real network, NAT from the simulated network to the real network, and forwarding of DNS queries to the real network.

There is also a convenience command named **connect-real-network** that automatically sets up NAT for outgoing traffic, forwarding of DNS queries to the real network, and incoming port forwarding for some common services. If there is no Ethernet link object, one is created and set up.

The **list-port-forwarding-setup** command describes the current port forwarding setup: it will list all incoming and outgoing ports, as well as the NAT and DNS forwarding status.

Note: Pinging between the simulated network and the real network will not work when using port forwarding, so ping should not be used to test if the connection is working. Ping uses the ICMP protocol, but only TCP and UDP traffic is supported with port forwarding.

The **connect-real-network** Command

The **connect-real-network** command is a convenience command that sets up NATP for outgoing traffic, enables forwarding of DNS queries to the real network, and opens incoming ports for FTP, HTTP and telnet to a simulated machine. This is an easy way to get inbound and outbound access for common services on a simulated machine.

The command requires a *target-ip* argument that specifies the IP address of the simulated machine that should be targeted by the incoming traffic. If there are multiple simulated machines, **connect-real-network** can be run once for each machine. Simics will select different ports on the simulation host for the incoming services for each simulated machine, and the selected ports are printed in the Simics console.

The **connect-real-network** command does not require an Ethernet link as argument, unless there is more than one in the simulation. If there is no Ethernet link or service node, they will be created automatically.

Example The **connect-real-network** allows us to set up all connections that are needed for most simple real network uses with one simple command. We can start from the checkpoint prepared in section 4.1, and then run the **connect-real-network** command with the IP address 10.10.0.40, which is the default address of *Firststeps*:

```
simics> connect-real-network 10.10.0.40
No ethernet link found, created default_eth_switch0.
Connected mpc8641d_simple.eth[0] to default_eth_switch0
Connected mpc8641d_simple.eth[1] to default_eth_switch0
Connected mpc8641d_simple.eth[2] to default_eth_switch0
Connected mpc8641d_simple.eth[3] to default_eth_switch0
Created instantiated 'std-service-node' component 'default_service_node0'
Connecting 'default_service_node0' to 'default_eth_switch0' as 10.10.0.1
NAPT enabled with gateway 10.10.0.1/24 on link default_eth_switch0.link.
NAPT enabled with gateway fe80::2220:20ff:fe20:2000/64 on link default_eth_switch0.li
Host TCP port 4021 -> 10.10.0.40:21
Host TCP port 4023 -> 10.10.0.40:23
Host TCP port 4080 -> 10.10.0.40:80
Real DNS enabled at 10.10.0.1/24 on link default_eth_switch0.link.
Real DNS enabled at fe80::2220:20ff:fe20:2000/64 on link default_eth_switch0.link.
```

The output shows that an **ethernet_switch** and a **std-service-node** components have been automatically created and connected to the simulated machine. NATP, DNS forwarding, and incoming port forwarding for FTP, HTTP and telnet have also been enabled.

Now start the simulation. Since we gave the service node the IP address 10.10.0.1, *Firststeps* should be configured with 10.10.0.1 as default gateway:

```
~ # route add default gw 10.10.0.1
```

It should now be possible to telnet from the simulated machine to hosts on the real network. In this case, we telnet to a Solaris machine with IP address 10.0.0.240; replace this address with any host answering to telnet on the network:

```
~ # telnet 10.0.0.240
Trying 10.0.0.240...
Connected to 10.0.0.240.
Escape character is '^]'.
```

```
SunOS 5.9
```

```
login: joe
Password:
Sun Microsystems Inc.      SunOS 5.9      Generic May 2002
$ exit
Connection closed by foreign host.
~ #
```

Firststeps can be configured to use the service node as DNS server and use it to look up real DNS names. To do that, add the line `nameserver 10.10.0.1` in the file `/etc/resolv.conf` on the simulated machine:

```
~ # echo nameserver 10.10.0.1 > /etc/resolv.conf
```

It should now be possible to look up the addresses of real hosts on the simulated machine, for example, `gnu.org`. *Firststeps* does not have the tools to perform DNS lookups. Instead, verify that DNS works by connecting to a real server by name:

```
~ # telnet gnu.org 80
GET /
<?xml version="1.0" encoding="ISO-8859-1"?>
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Strict//EN"
    "http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd">
<html xmlns="http://www.w3.org/1999/xhtml" lang="en" xml:lang="en">
<head>
<title>Server error!</title>
[...]
```

Connection closed by foreign host.

```
~ #
```

FTP, HTTP and telnet servers running on the simulated machine should also be accessible. *Firststeps* runs both a telnet and a HTTP server. Just use port 4023 and 4080 instead of 23 and 80. The exact ports of the host these services are mapped to varies if the default ports are already in use. Look at the output from the **connect-real-network** above for the port numbers to use.

Incoming Port Forwarding

The **connect-real-network-port-in** command sets up port forwarding from a port on the host machine to a specific port on a simulated machine. It takes three required arguments: *ethernet-link*, *target-ip* and *target-port*, that specify the Ethernet link, IP address and port the traffic should be forwarded to.

An IP address and preferred port can be selected for incoming traffic on the simulation host using the *host-ip* and *host-port* arguments. If these arguments are not provided, Simics will select a port automatically and print it on the Simics console, and receive all IPv4 traffic (i.e., IP 0.0.0.0) from that port. In order to forward multicast traffic, specify that multicast address (e.g., specify 239.255.255.253 to forward IPv4 SLP traffic).

The **connect-real-network-port-in** command can also take the flags *-tcp* and *-udp*, which specify whether forwarding is set up for a TCP or a UDP port. If neither is provided, forwarding will be set up for both the TCP and UDP ports.

The service node acts as a proxy for incoming traffic, so to initiate a connection to a specific port on the simulated machine, the real machine should contact the corresponding open port on the simulation host. The simulation host is not a gateway to the simulated network.

Any UDP packets sent to a port on the simulation host are forwarded to the specified port and IP address on the simulated network. For the simulated machine to be able to return UDP packets to the real network, a separate forwarding rule must be set up using the **connect-real-network-port-out** command.

Any TCP connections to the port on the simulation host are forwarded to the specified port and IP address on the simulated network. Since TCP connections are two-ways, once a connection has been established, data can be sent in both directions.

Note: You will probably have to manually add the incoming ports to your host based software firewall if you want to access the simulated network from another machine. Tracking down network problems when you forget to update the firewall is annoying as the packets tend to get dropped silently without a log.

The FTP protocol needs to open additional ports when transferring files. Simics handles this by automatically opening outgoing ports for FTP when needed, so FTP will work as long as it is in *active* mode.

Example *Firststeps* runs telnet on port 23. We can now set up a port forwarding rule that allows us to access the telnet service from the real network. Start from the checkpoint, create an Ethernet link and service node, connect the simulated machine to the Ethernet link and run the **connect-real-network-port-in** command like this:


```

simics> load-module eth-links
simics> new-ethernet-switch switch0
Created instantiated 'ethernet_switch' component 'switch0'
simics> new-std-service-node sn0
Created instantiated 'std-service-node' component 'sn0'
simics> sn0.connect-to-link switch0 10.10.0.1
Adding host info for IP 10.10.0.1: simics0.network.sim  MAC: 20:20:20:20:20:00
simics> connect mpc8641d_simple.eth[0] switch0.device1
simics> connect-real-network-port-in ethernet-link = switch0
target-ip = 10.10.0.40 target-port = 23 host-port = 2023 -tcp
Host TCP port 2023 -> 10.10.0.40:23
simics> enable-real-time-mode

```

Firststeps uses the IP address 10.10.0.40 and the telnet service runs on TCP port 23. We use port 2023 on the simulation host, but any free port can be used. The last command slows down *Firststeps* to avoid time-outs.

Start the simulation, then start a telnet from a real host to the telnet port of the simulated machine by connecting to port 2023 of the simulation host. In our case, we do the telnet on the simulation host itself; replace *localhost* with the IP address of the simulation host if running telnet from another machine:

```

computer$ telnet -l root localhost 2023
Trying 127.0.0.1...
Connected to localhost.
Escape character is '^]'.

(none) login: root
~ # exit
Connection closed by foreign host.
computer$

```

Outgoing Port Forwarding

The **connect-real-network-port-out** command sets up port forwarding from a port on a service node to a specific port on a host on the real network. It takes four required arguments: *service-node-port* *ethernet-link*, *target-ip* and *target-port*, that specify the port on the **service node** that will forward traffic to the target, the Ethernet link the service node is connected to, and the real IP address and port to which the traffic should be forwarded.

The command can optionally take the flags *-tcp* and *-udp*, to specify whether the forwarding should be set up for a TCP or UDP port. If neither is provided, forwarding will be set up for both the TCP and UDP port.

The service node acts as a proxy for outgoing traffic, so to initiate a connection to a port on a host on the real network, the simulated machine should connect to the corresponding mapped port on the service node. The service node is not a gateway to the real network.

Any UDP packets sent to a port on the service node are forwarded to the specified port and IP address on the real network. For the real host to be able to return UDP packets to the simulated network, a separate forwarding rule must be set up using the **connect-real-network-port-in** command.

Any TCP connections to the port on the service node are forwarded to the specified port and IP address on the real network. Since TCP connections are two-ways, once a connection has been established data can be sent in both directions.

Example By setting up forwarding from a port on a service node to port 23 of a host on the real network, it should be possible to telnet to the real host by connecting to the port on the service node from *Firststeps*. We can start from the checkpoint we prepared in section 4.1, and create an Ethernet link and a service node, connect the simulated machine to the Ethernet link and run the **connect-real-network-port-out** command. Here we use a host on the real network with IP address 10.0.0.240, replace it with the IP address of a real host on the network:

```
simics> load-module eth-links
simics> new-ethernet-switch switch0
Created instantiated 'ethernet_switch' component 'switch0'
simics> new-std-service-node sn0
Created instantiated 'std-service-node' component 'sn0'
simics> sn0.connect-to-link switch0 10.10.0.1
Adding host info for IP 10.10.0.1: simics0.network.sim MAC: 20:20:20:20:20:00
simics> connect mpc8641d_simple.eth[0] switch0.device1
simics> connect-real-network-port-out service-node-port = 2323 ↵
ethernet-link = switch0 target-ip = 10.0.0.240 target-port = 23 -tcp
10.10.0.1 TCP port 2323 on link switch0 -> host 10.0.0.240:23
simics> enable-real-time-mode
```

Now start the simulation. We used the IP address 10.10.0.1 and the port 2323 for the service node, so we should be able to telnet to the real host by connecting to port 2323 of 10.10.0.1 from *Firststeps*:

```
~ # telnet 10.10.0.1 2323
Trying 10.10.0.1...
Connected to 10.10.0.1.
Escape character is '^]'.
```

```
SunOS 5.9
```

```
login: joe
Password:
No directory! Logging in with home=/
Last login: Sun Jun 2 07:45:58 from 10.0.0.211
```

```
Sun Microsystems Inc.   SunOS 5.9           Generic May 2002
$ exit
~ #
```

NAPT

The **connect-real-network-napt** command sets up *NAPT* (*network address port translation*, also known as just *NAT* or *network address translation*) between the simulated network and the real network. With NAPT enabled, the service node will act as a gateway on the simulated network and automatically mediate TCP connections to the real network.

The **connect-real-network-napt** only has one required argument, *ethernet-link*, that specifies the Ethernet link that should be connected to the real network.

The simulated machines must be configured to use the service node as gateway for the real network, so that it is able to capture the outgoing traffic. The simulated machines will then be able to access hosts on the real network using their real IP addresses. By combining NAPT with DNS forwarding, described in section 4.2.1, the real DNS names of hosts on the real network can be used as well.

The NAPT setup is not specific to a simulated machine, so **connect-real-network-napt** needs only to run once for each Ethernet link, and all simulated machines on the link get outbound access.

Since NAPT only allows new TCP connections to be opened from the simulated network to the real network, and the FTP protocol need to open new ports when transferring files, *passive* mode FTP should be used when connecting to an FTP server on a host on the real network from a simulated machine. An alternative is to use the FTP server implemented in the Simics service-node and avoid the need to connect to a real network.

Example To try NAPT, we can start from the checkpoint we prepared in section 4.1, create an Ethernet link and service node, connect the simulated machine to the Ethernet link and run the **connect-real-network-napt** command like this:

```
simics> load-module eth-links
simics> new-ethernet-switch switch0
Created instantiated 'ethernet_switch' component 'switch0'
simics> new-std-service-node sn0
Created instantiated 'std-service-node' component 'sn0'
simics> sn0.connect-to-link switch0 10.10.0.1
Adding host info for IP 10.10.0.1: simics0.network.sim MAC: 20:20:20:20:20:00
simics> connect mpc8641d_simple.eth[0] switch0.device1
simics> connect-real-network-napt ethernet-link = switch0
NAPT enabled with gateway 10.10.0.1 on link switch0.link
NAPT enabled with gateway fe80::2220:20ff:fe20:2000/16 on link switch0.link
simics> enable-real-time-mode
```

The simulated machine should be configured to use the service node as its default gateway:

```
~ # route add default gw 10.10.0.1
```

DNS Forwarding

The **enable-real-dns** and **disable-real-dns** commands of the service node enable and disable forwarding of DNS requests to the real network by a service node. This allows simulated machines to look up names and IP addresses of hosts on the real network, using the service node as DNS server.

Example To try DNS forwarding, we can start from the checkpoint we prepared in section 4.1, and create an Ethernet link and a service node, connect the simulated machine to the Ethernet link and run the **enable-real-dns** command like this:

```
simics> load-module eth-links
simics> new-ethernet-switch switch0
Created instantiated 'ethernet_switch' component 'switch0'
simics> new-std-service-node sn0
Created instantiated 'std-service-node' component 'sn0'
simics> sn0.connect-to-link switch0 10.10.0.1
Adding host info for IP 10.10.0.1: simics0.network.sim MAC: 20:20:20:20:20:00
simics> connect mpc8641d_simple.eth[0] switch0.device1
simics> sn0.enable-real-dns
Host TCP port 2023 -> 10.10.0.40:23
simics> enable-real-time-mode
```

To tell *Firststeps* to use the service node as DNS server, the line `nameserver 10.10.0.1` is needed in the file `/etc/resolv.conf`:

```
~ # echo nameserver 10.10.0.1 > /etc/resolv.conf
```

4.2.2 Ethernet Bridging

Simics can act as a bridge between simulated Ethernet networks and the real Ethernet networks of the host. With this type of connection, the simulated machines will appear as directly connected to the real network, both to the simulated machines and to the hosts on the real network. For that reason, the simulated machines should be configured with IP addresses from the same subnet as the real hosts.

Since the simulated machines appear to be located on the real network, there is no need to configure routes on real hosts that communicate with it. They can find the simulated machines by sending ARP requests, just like they would find other real hosts.

When using a bridged connection, it is recommended to use a dedicated host network interface for the bridge, not the interface that the host uses for general network access. Otherwise, special configuration of the bridge interface is necessary for any kind of network

access for the host. The details about this configuration will not be covered in this manual. See section 4.3 for how to set up TAP access.

To create a bridged connection to the real network, use the **connect-real-network-bridge** command. It takes an *interface* argument that specifies which host Ethernet interface should be used. It also takes an *host-access* argument that tells Simics how to access the host's Ethernet interface; by default, this is `TAP`, but the deprecated `raw` can also be chosen. TAP access requires some additional setup, described after the example below.

By default Simics will translate MAC addresses between the simulated network and the real network, so that the host's MAC address is used for all packets sent on the real network. This has two advantages: the simulation host does not have to listen on the real network in promiscuous mode, reducing its load; it also avoids problems with MAC address collisions between multiple simulations connected to the same real network. However, each simulated machine communicating with the real network must still be configured with a unique IP address.

There are a couple of drawbacks with MAC address translation. One is that it is limited to ARP and IPv4 packets; other kinds of packets will not be bridged. There may also be problems with setting up routing on simulated machines. Since the destination MAC addresses are translated, the packets may not go where expected.

To be able to bridge other kinds of traffic than ARP and IPv4, for example, DHCP, IPv6, or IPX, MAC address translation can be turned off by specifying the *-no-mac-xlate* flag to **connect-real-network-bridge**. This enables all kinds of Ethernet traffic to be bridged between the simulated network and the real network. Care must be taken that each simulated machine connected to the real network is configured with a unique MAC address, to avoid MAC address collisions on the real network.

A big drawback with disabling MAC address translation is that the host must listen on the real network in promiscuous mode. This increases the risk of dropping packets from the real network that were intended for the simulated network, because unrelated traffic will be bridged to the simulated network. It is therefore recommended that bridging without MAC address translation only be used on dedicated networks, or on network interfaces with very little unrelated traffic.

Example

This example assumes that the simulation is starting from the checkpoint prepared in section 4.1, and that a TAP bridge has been set up as described in section 4.2.2.

To set up an Ethernet bridging connection between the real network and the simulated network, run the **connect-real-network-bridge** command. This will automatically create an Ethernet link, connect it to the simulated machine and set up bridging to the real network:

```
simics> connect-real-network-bridge
No ethernet link found, created default_eth_switch0.
Connected mpc8641d_simple.eth[0] to default_eth_switch0
Connected mpc8641d_simple.eth[1] to default_eth_switch0
Connected mpc8641d_simple.eth[2] to default_eth_switch0
Connected mpc8641d_simple.eth[3] to default_eth_switch0
TAP does not support MAC address translation; turning it off
```

```
[rn0.rn info] Connecting to existing TAP device 'sim_tap0'
'default_eth_switch0' connected to real network.
Created 'real-network-bridge' component 'rn0'
```

When using Ethernet bridging, the simulated machine should be configured with an unused IP address and netmask from the real network. In this case we use 10.0.0.241 and 255.255.255.0. Replace it with an unused IP address and netmask from the real network:

```
~ # ifconfig eth0 10.0.0.241 netmask 255.255.255.0
```

The simulated machine is now connected to the real network. Any kind of IP traffic is bridged between the simulated network and the real network. It should be possible to ping any real host from the simulated machine. Replace *10.0.0.240* with the address of a host on the real network:

```
~ # ping 10.0.0.240 -c 3
PING 10.0.0.240 (10.0.0.240): 56 data bytes
64 bytes from 10.0.0.240: seq=0 ttl=64 time=10.285 ms
64 bytes from 10.0.0.240: seq=1 ttl=64 time=0.349 ms
64 bytes from 10.0.0.240: seq=2 ttl=64 time=0.323 ms

--- 10.0.0.240 ping statistics ---
3 packets transmitted, 3 packets received, 0% packet loss
round-trip min/avg/max = 0.323/3.652/10.285 ms
```

Of course, it should also be possible to ping the simulated machine from the real host.

Running `tracert` on the simulated machine shows that it is connected directly to the real network; there are no routers between it and the real host. Again, replace *10.0.0.240* with a host on the real network:

```
~ # tracert 10.0.0.240
tracert to 10.0.0.240 (10.0.0.240), 30 hops max, 38 byte packets
 1  10.0.0.240 (10.0.0.240)  10.106 ms  0.371 ms  0.347 ms
```

If the IP address of the simulated machine itself is printed and `!H` is printed after the response times, it means the simulated machine can not reach the real host, and the configuration is incorrect.

Setting up TAP for Ethernet Bridging

To use TAP access with the `connect-real-network-bridge` command, the operating system must be configured to act as a bridge between the virtual interface and the real interface. Follow the steps below to set it up.

Note: When setting up bridging between a TAP interface and a real Ethernet interface, the host will no longer accept traffic on the real interface. All connections the host has open on the interface will be lost. We therefore strongly recommend that bridging is only set up on dedicated host interfaces.

Unix:

1. Create a TAP interface, as described in section [4.3.1](#)
2. Create a bridge interface and connect the TAP interface and the real interface. Turn off STP (Spanning Tree Protocol) in the bridge as well, otherwise there will be STP traffic from the bridge into both the simulated and the real network. Here the name of the created bridge interface is `sim_br0` and the interface used is `eth1`, but other names and interfaces can be used.

```
computer# brctl addbr sim_br0
computer# brctl addif sim_br0 sim_tap0
computer# brctl addif sim_br0 eth1
computer# brctl stp sim_br0 off
```

3. Bring up the TAP interface and the bridge interface.

```
computer# ifconfig sim_tap0 promisc up
computer# ifconfig sim_br0 promisc up
```

4. Bring up the Ethernet interface.

```
computer# ifconfig eth1 up
```

To remove the bridging when finished, do the following:

1. Bring down the TAP interface and the bridge interface.

```
computer# ifconfig sim_tap0 down
computer# ifconfig sim_br0 down
```

2. Delete the bridge interface.

```
computer# brctl delbr sim_br0
```

Note: The `brctl` utility is usually not present in default Linux installations. It is usually included in the `bridge-utils` package.

Windows:

1. Open the *Network Connections* folder of *Control Panel*.
2. Select the TAP-Win32 interface and the real interface to bridge to at the same time.
3. Right-click on the TAP-Win32 interface to bring up a context menu, and select *Bridge Connections*.

Windows will now set up bridging between the TAP-Win32 interface and the real interface. If successful both the TAP-Win32 and real interfaces will seem to disappear as they are now contained in the new bridge interface. To undo the bridging, go to the *Network Connections* folder, right-click on the bridge and select *Delete*.

4.2.3 Host Connection

Simics can connect a simulated network to a virtual Ethernet (TAP) interface on the simulation host. The simulation host and the simulated machines will then be able to communicate as if they were connected to the same Ethernet network. For example, by configuring the simulation host with an IP address on the TAP interface, the simulation host and the simulated machines will be able to send IP traffic to each other.

Enabling IP forwarding on the host will also allow the simulated machines to access other hosts on the real network, using the host operating system's IP routing facilities. Read the instructions after the example below for instructions about how to set it up.

To connect the simulated network to the TAP interface, the TAP interface should be configured on the simulation host, as described in section 4.3.1. Use the **connect-real-network-host** command, which simply takes the name of the TAP interface as the *interface* argument. The simulation host will now appear on the simulated network. Configure the TAP interface of the host with an IP address from the same subnet as the simulated machines, and the simulated machines will be able to communicate with the host.

Simulated machine configurations provided with Simics usually use IP addresses from the 10.10.0.x subnet, so the simulation host should typically get an IP address on the form 10.10.0.x with a netmask of 255.255.255.0.

On Unix, this is configured with `ifconfig`, which requires administrative privileges:

```
computer# ifconfig sim_tap0 10.10.0.x netmask 255.255.255.0 up
```

On Windows, use these steps instead:

1. Open the *Network Connections* folder of the *Control Panel*.
2. Right-click on the TAP-Win32 interface to bring up a context menu, and select *Properties*.
3. A property dialog box will open. Select *Internet Protocol (TCP/IP)* and click on the *Properties* button.

4. A new property dialog box will open. Select *Use the following IP address* and enter the IP address and Subnet mask the simulation host should use on the simulated network.
Select an IP address on the form 10.10.0.x and the netmask 255.255.255.0. The *Default gateway* field can typically be left blank.
5. Click OK in both property dialog boxes.

Example

This example assumes that the simulation is starting from the checkpoint prepared in section 4.1, that there is a correctly set up TAP interface on the simulation host for host connection according to section 4.3.1, and that it has been configured with the IP address 10.10.0.1 as described above. Here the name if the TAP interface is assumed to be `sim_tap0`: replace it with the name of the TAP interface.

To connect the TAP interface to the simulated network, use the **connect-real-network-host** command:

```
simics> connect-real-network-host interface = sim_tap0
No ethernet link found, created default_eth_switch0.
Connected mpc8641d_simple.eth[0] to default_eth_switch0
Connected mpc8641d_simple.eth[1] to default_eth_switch0
Connected mpc8641d_simple.eth[2] to default_eth_switch0
Connected mpc8641d_simple.eth[3] to default_eth_switch0
[rn0.rn info] Receive buffer size 262142 bytes.
[rn0.rn info] Using mmap packet capture.
[real_net0 info] Connecting to existing TUN/TAP device 'sim_tap0'
'default_eth_switch0' connected to real network.
```

Note: On Windows, the message “Connecting to existing TUN/TAP device ‘sim_tap0’” will not appear.

Any kind of Ethernet traffic can now pass between the simulated network and the simulation host. It should be possible, for example, to ping the simulation host from the simulated machine, at the IP address configured on the TAP interface:

```
~ # ping 10.10.0.1 -c 5
PING 10.10.0.1 (10.10.0.1) from 10.10.0.40 : 56(84) bytes of data.
64 bytes from 10.10.0.1: icmp_seq=1 ttl=64 time=1.15 ms
64 bytes from 10.10.0.1: icmp_seq=2 ttl=64 time=1.11 ms
64 bytes from 10.10.0.1: icmp_seq=3 ttl=64 time=10.9 ms
64 bytes from 10.10.0.1: icmp_seq=4 ttl=64 time=1.11 ms
64 bytes from 10.10.0.1: icmp_seq=5 ttl=64 time=1.11 ms

--- 10.10.0.1 ping statistics ---
5 packets transmitted, 5 received, 0% loss, time 4037ms
```

```
rtt min/avg/max/mdev = 1.113/3.085/10.932/3.923 ms
```

IP Forwarding

Enabling IP forwarding on the simulation host will allow real machines to access the simulated network by routing the traffic through the simulation host. This method is called *IP forwarding*.

On Unix run the following command to set up ip forwarding:

```
computer# sysctl -w net.ipv4.ip_forward=1  
net.ipv4.ip_forward = 1
```

And to disable IP forwarding again:

```
computer# sysctl -w net.ipv4.ip_forward=0  
net.ipv4.ip_forward = 0
```

To set it up on Windows, change the registry key `HKEY_LOCAL_MACHINE\SYSTEM\CurrentControlSet\Services\Tcpip\Parameters\IPEnableRouter` from 0 to 1 and reboot. If the key does not exist, create it and give it the value 1. To disable IP forwarding, reset the registry key to 0 and reboot.

For the routing to work, both the simulated machines and the machines on the real network must be configured: on the simulated machines, a route must be added to the real network with the IP address of the host's TAP interface as gateway; on the real machines, it is a route to the simulated network with the IP address of the simulation host's Ethernet interface that should be added.

4.3 Accessing Host Ethernet Interfaces

When connecting to a real network using other connection types than port forwarding, Simics needs low-level access to the simulation host's Ethernet interfaces to send and receive packets. However, operating systems do not usually allow user programs low-level access to Ethernet interfaces: it requires specific configuration with administrative rights.

The access method used by Simics is called *TAP*; it is a virtual Ethernet interface on the simulation host. Ethernet bridging and host connections require the use of TAP, but, as stated above, port forwarding connections do not.

4.3.1 TAP Access

With TAP access Simics will connect the simulated network to a virtual Ethernet (TAP) interface provided by the operating system. Accessing the TAP interface does not require administrative privileges, so once the TAP interface has been configured, Simics can connect to the real network without administrative privileges.

The TAP interface can either be bridged to a real Ethernet interface to create an Ethernet bridging connection, or configured with an IP address to create a host connection. This section only describes the basic configuration which is required for both kinds of connections. Read section 4.2.2 for the additional steps needed for Ethernet bridging connections and 4.2.3 for the additional steps required for host connections.

The setup required for TAP access differs for simulation hosts running Unix and Windows.

Unix

Creating a TAP interface that Simics can use is done in two simple steps. These commands require administrative privileges:

1. Give the user running the simulation access to the `/dev/net/tun` device.

```
computer# chmod 666 /dev/net/tun
```

2. Create the TAP interface. Here the name of the user that will be using the TAP interface is assumed to be `joe` and the name of the TAP interface will be `sim_tap0`, but it should of course be replaced with the correct user name and TAP interface.

```
computer# tunctl -u joe -t sim_tap0  
Set 'sim_tap0' persistent and owned by uid 4711
```

To remove a TAP interface, use the `-d` argument to `tunctl`. Again, this requires administrative privileges:

```
computer# tunctl -d sim_tap0  
Set 'sim_tap0' nonpersistent
```

Note: The `tunctl` utility is usually not present in default Linux installations. It is usually included in the User Mode Linux package.

Windows

To create a TAP interface, the driver called *TAP-Win32* is required, that is part of the *OpenVPN* package. See the *Installation Guide* for information about how to do this. Administrative privileges are required on the simulation host both to install *TAP-Win32* and to connect Simics to the real network.

One *TAP-Win32* interface is created automatically when installing *TAP-Win32*. However, to have more than one connection to the real network at a time, multiple *TAP-Win32* interfaces are necessary.

To create additional interfaces, run the script `addtap.bat` in the `bin` directory in the *OpenVPN* installation directory. Typically, the path to the script will be `C:\Program`

`Files\OpenVPN\bin\addtap.bat`. When Windows warns that the device driver has not passed compatibility testing, just select *Continue Anyway*. When done, there should be a new TAP-Win32 interface in the *Network Connections* folder of the *Control Panel*.

To remove TAP-Win32 interfaces, run the `deltapall.bat` script, located in the same directory as the `addtap.bat` script. Note that this will remove all TAP-Win32 interfaces.

Note: Simics uses the MAC addresses of the simulation host's network interfaces to identify the host for licensing purposes. The licensing system does not differentiate TAP-Win32 interfaces and real Ethernet interfaces. Creating or removing TAP-Win32 interfaces may therefore cause a node-locked license to stop working. Renaming the TAP device so that the Local Area Connection is first will probably fix this.

Configuring a TAP interface consists of setting its name and IP address.

1. select Start Menu → Settings → Network Connections
2. rename the new TAP interface (name is likely to similar to "TAP-Win32 Adapter V8") easily identifiable like "simicsTAP"
3. to configure the IP address, right click the TAP device and choose Properties, then select "Internet Protocol(TCP/IP)" and click the "Properties" button. Some values are 10.10.0.100 for the IP address and 255.255.255.0 for the Subnet Mask.
4. Click "OK" twice to save the new configuration

TAP devices with multiple users and simulations

On Windows, TAP devices are typically created with sufficient permissions for all users to be able to access them. On Linux however, you have to be the user or belong to the group that was specified when the device was created to be able to access it . On both Windows and Linux, each device can only be opened by one application at a time. These facts can lead to problems when multiple users tries to run the same Simics configuration on the same machine, possibly even at the same time.

Let us take a closer look at how to solve the problem with having multiple users running identical configurations on the same machine, but not necessarily at the same time:

Newer versions of the Linux kernel offers the possibility to assign a group to a TAP device; any member of that group is allowed to open the tap device, and no further configuration is needed. That means that any routing table entries that has been set up on the simulation host has to be configured only once and can remain the same after that. It is necessary to have a version of `tunctl` recent enough to take advantage of this possibility. Without the group functionality of TAP devices, each user must use its own TAP device. This means that the interface argument to **connect-real-network-[host,bridge]** must depend on which user that runs the command, and routes on the simulation host must be updated to match that TAP device. As has previously been said, TAP devices on Windows can typically be accessed by all users, so the choice of having one or multiple TAP devices is entirely up to the users.

If several users needs to run identical configurations at the same time with access to the same real network it gets even more complicated; the simple answer is that it is generally not

possible. Even if the users run the configuration on different simulation hosts, it is generally not possible connect both simulations to the real network in a correct and working manner. The reasons are a little different between a host connection and a bridged connection, but in both cases it has to do with the fact that there are multiple target machines that have the same IP and MAC addresses on connected to the real network in some way.

When using a bridged real network connection, the simulated network will appear to be the same network as the simulated network. When running two identical simulations, both MAC and IP addresses will be duplicated on the same network, which will lead to unexpected behavior even if different simulation hosts are used. Even with MAC address translation the IP addresses will still collide. When using a host connection the simulation host can only route traffic to one of the simulated networks since the network addresses and TAP interface addresses has to be the same for both simulations. The only possible setup is a host connection where each simulation is running on a different machine, and the only communication is internally on the simulated network and between simulated machines and their simulation host.

The general solution to this problem is to use configurations that are not completely identical. Each simultaneously running configuration should have unique network properties, such as IP address range, MAC address range and TAP interface. Using such Simics configurations, all network configuration, including routes and TAP interfaces on the real hosts can be done in advance and does not have to be changed for every Simics configuration that is started. For example, if a bridged connection is required, several TAP interfaces and a real network interface can be bridged together to form a single bridge.

The examples below show how real network connections can be configured when running multiple *almost* identical Simics configurations simultaneously. For both examples, the part of the configuration that has to be done by the root user can be done once so that the user running the simulation does not need administrative privileges.

Figure 4.2 shows a setup with two Simics processes on the same simulation host, both using a host connection. Each Simics process connects to its own TAP device. All target machines in the `Simics 0` configuration has IP addresses on the 10.10.x.0 network while the `Simics 1` configuration has IP addresses on the 10.10.y.0 network. It is important that all the real hosts that the target machines needs to talk to has static routes set up for the 10.10.x.0 and 10.10.y.0 networks, using the simulation host as gateway.

Figure 4.3 shows a setup similar to Figure 4.2. In this case the target machine connect to the real network using an Ethernet bridge on the simulation host. The bridge is invisible to both the target machines and the real hosts. As all IP addresses must be on the 192.168.0.0 network, the `Simics 0` configuration uses IP addresses between 192.168.0.0 and 192.168.0.9, and the `Simics 1` configuration uses IP addresses between 192.168.0.10 and 192.168.0.19. Note that both `sim_tap0` and `sim_tap1` is part of the same Ethernet bridge, and that none of the interfaces that is a part of the bridge has any IP address. In this description the MAC addresses has been left out to not unnecessarily clutter the figure and description. These needs to be partitioned between the two Simics configurations in the same way as the IP addresses.

4.3. Accessing Host Ethernet Interfaces

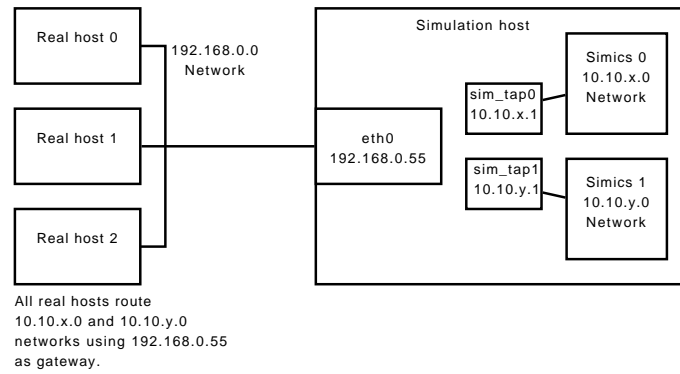


Figure 4.2: Example setup of multiple Simics instances using a host connection




Figure 4.3: Example setup of multiple Simics instances using a bridged connection

4.4 Selecting Host Ethernet Interface

When connecting to a real network on a host with multiple network interfaces installed, Simics will select one of them for the real network connection. If the default selection is incorrect, all commands have an *interface* argument to select the desired interface. The interface name expected by the *interface* argument is the ordinary interface name used by the host operating system. The next two sections describe how to obtain the names of the simulation host's network interfaces for Unix and Windows hosts.

4.4.1 Unix

All network interfaces are listed by `/sbin/ifconfig -a` on the simulation host:

```
computer$ /sbin/ifconfig -a
eth0      Link encap:Ethernet  HWaddr 00:10:18:0A:DE:EF
          BROADCAST MULTICAST  MTU:1500  Metric:1
          RX packets:0 errors:0 dropped:0 overruns:0 frame:0
          TX packets:0 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:1000
          RX bytes:0 (0.0 b)  TX bytes:0 (0.0 b)
          Interrupt:21

eth1      Link encap:Ethernet  HWaddr 00:0C:F1:D1:FF:09
          inet addr:10.0.0.140  Bcast:10.0.0.255  Mask:255.255.255.0
          UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
          RX packets:671467287 errors:0 dropped:0 overruns:0 frame:0
          TX packets:647635204 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:1000
          RX bytes:3725791210 (3553.1 Mb)  TX bytes:217046405 (206.9 Mb)
          Interrupt:20 Base address:0xdf40 Memory:fceef000-fceef038

lo        Link encap:Local Loopback
          inet addr:127.0.0.1  Mask:255.0.0.0
          UP LOOPBACK RUNNING  MTU:16436  Metric:1
          RX packets:24929 errors:0 dropped:0 overruns:0 frame:0
          TX packets:24929 errors:0 dropped:0 overruns:0 carrier:0
          collisions:0 txqueuelen:0
          RX bytes:4164218 (3.9 Mb)  TX bytes:4164218 (3.9 Mb)
```

For example, to use the first interface listed above, specify `eth0` as the *interface* argument.

4.4.2 Windows

The interface name expected by the *interface* argument is the ordinary interface name used by the host operating system. All interfaces of the host are listed in the *Network Connections*

folder in the *Control Panel*; alternatively, `ipconfig` can be used in a command prompt on the simulation host:

```
C:\>ipconfig
```

```
Windows IP Configuration
```

```
Ethernet adapter Local Area Connection:
```

```

Connection-specific DNS Suffix  . :
IP Address. . . . . : 10.0.0.191
Subnet Mask . . . . . : 255.255.255.0
Default Gateway . . . . . : 10.0.0.3

```

```
Ethernet adapter LAB:
```

```

Connection-specific DNS Suffix  . :
IP Address. . . . . : 10.10.0.1
Subnet Mask . . . . . : 255.255.255.0
Default Gateway . . . . . :

```

```
Ethernet adapter TAP:
```

```
Media State . . . . . : Media disconnected
```

For example, to use the first interface listed above, specify “*Local Area Connection*” as the *interface* argument. There is no need to specify the entire interface name: it is enough to specify a unique substring.

4.5 Performance

When using other connection types than port forwarding, Simics has to prevent the simulated network from being flooded with packets from the real network. In practice, the simulated machine is not always able to handle the amount of traffic it receives, and Simics cannot buffer all incoming packets. This is especially a problem when the simulation is stopped.

Flooding is prevented by limiting the amount of traffic that is allowed to enter the simulated network per simulated second. This amount is determined by the *tx_bandwidth* and *tx_packet_rate* attributes of the real network object created for the connection, typically named **real_net0**. The unit of the *tx_bandwidth* attribute is bits per simulated second, and the unit of the *tx_packet_rate* attribute is packets per simulated second. Either can be set to unlimited by setting them to 0, but this is not recommended unless a very limited amount of data is to be received. The default is to allow 10 megabits per simulated second and an unlimited number of packets.

In addition to allowing the selected rate of traffic into the simulated network, Simics buffers traffic for an additional 0.1 seconds of simulated time, to avoid dropping packets during a short peak in the traffic. If more packets arrive once this buffer is full, Simics will drop them.

To get better performance out of the connection to the real network, modify the value of the `tx_bandwidth` and `tx_packet_rate` attributes until satisfied. A good strategy is to set the `tx_bandwidth` attribute to the amount of traffic that the simulated machine is expected to handle per simulated second, and then try to increase the `tx_packet_rate` from about 5000 and see at what packet rate the best performance is achieved.

For example, this will set the limit to 100 megabits and 10000 packets per simulated second:

```
simics> @conf.real_net0.tx_bandwidth = 100000000
simics> @conf.real_net0.tx_packet_rate = 10000
```

4.6 Troubleshooting

A network monitoring tool such as *Wireshark* (formerly known as *Ethereal*) is invaluable when debugging problems with the real network connections. It is a graphical traffic analyzer that can analyze most common network protocols. Wireshark is available from <http://www.wireshark.com>.

There are some pitfalls one might encounter when trying to connect a simulated network to a real one:

Trying to access the simulation host

Accessing the simulation host from the simulated network, or the other way around, is only supported with port forwarding and host connection. It is difficult—sometimes impossible—to access the simulation host from the simulated network when setting up an Ethernet bridging connection.

Simulated OS has no route

When using a NAPT connection and the operating system of the simulated machine does not have a correct route to the real network, the simulated machine will drop the packets or send them to the wrong address. To view the routing setup on the simulated machine, use the command `netstat -r` on Linux or `route print` on Windows. Note that these commands should be executed on the *simulated* machines. The simulated OS should have a default route to the service node.

Simics uses the wrong host network interface

On a host with multiple network interfaces installed, Simics will only use one of them for a real network connection. If the default selection is incorrect, use the *interface* argument of the connect command to select the desired network interface. See the *Selecting Host Interface* part of section 4.4.

Index

B

BOOTP, [16](#)
brctl, [39](#), [40](#)
bridging, [27](#), [36](#)
 TAP, [27](#), [28](#)

C

connect-real-network, [29](#), [30](#)
 connect-real-network-port-in, [32](#)
 connect-real-network-port-out, [32](#), [33](#)
connect-real-network-, [28](#)

D

DHCP, [16](#)
 add-host, [16](#)
 dhcp-add-pool, [16](#)
DNS, [16](#), [29](#), [36](#)
 add-host, [16](#)

E

enable-real-time-mode, [7](#), [25](#)
eth-links, [10](#)
Ethernet, [10](#)
 bridging, [27](#), [28](#), [36](#)
Ethernet link, [10](#)
 ethernet-link, [29](#)
ethernet-cable, [10](#)
ethernet-hub, [10](#)
ethernet-vlan-switch, [11](#)
ethernet_switch, [10](#)

F

FTP, [28](#)
 active mode, [32](#)

G

gateway, [15](#)

H

host connection, [27](#), [28](#)
HTTP, [28](#)

I

ICMP, [27](#), [30](#)
IEEE 802.1Q, [11](#)
IP
 address, [15](#), [16](#)
 forwarding, [27](#), [42](#)
 routing, [15](#)
IPv4, [14](#)
IPv6, [14](#)

L

latency, [12](#)

M

MAC address, [16](#)
 no-mac-xlate, [37](#)
 translation, [37](#)
minimum latency, [13](#)

N

NAPT, [29](#), [30](#), [35](#)
NAT, [35](#)

P

Ping, [30](#)
port forwarding, [27–29](#)
 connect-real-network, [29](#)
 incoming, [29](#), [32](#)
 list-port-forwarding-setup, [29](#)
 outgoing, [29](#), [33](#)
port forwarding
 NAPT, [29](#)

R

routing, [15](#)

gateway, [15](#)
route-add, [15](#)

S

service node, [13](#), [29](#)
std-service-node, [14](#)

T

TAP, [37](#), [42](#)
TCP/IP, [10](#), [29](#)
telnet, [25](#), [28](#), [29](#)
TFTP, [17](#)
tunctl, [43](#)

U

UDP, [29](#)

V

VLAN, [11](#)